




*Empirical. 295 tool calls aggregated across 2 LLM(s) (DeepSeek-V3 + Qwen2.5-72B-Instruct) producing tool calls in response to adversarial prompts across 6 attack categories. Each call is fed to every protocol's host-side validator (the real dcp.bridge.Bridge for DCP; jsonschema for MCP / IoT-MCP / OpenAPI).
Corpus in tools/gen_llm_corpus.py + llm_corpus.json; aggregation in tools/bench_hallucination_empirical.py.*